

# Model Checking

## ESS 575 Models for Ecological Data

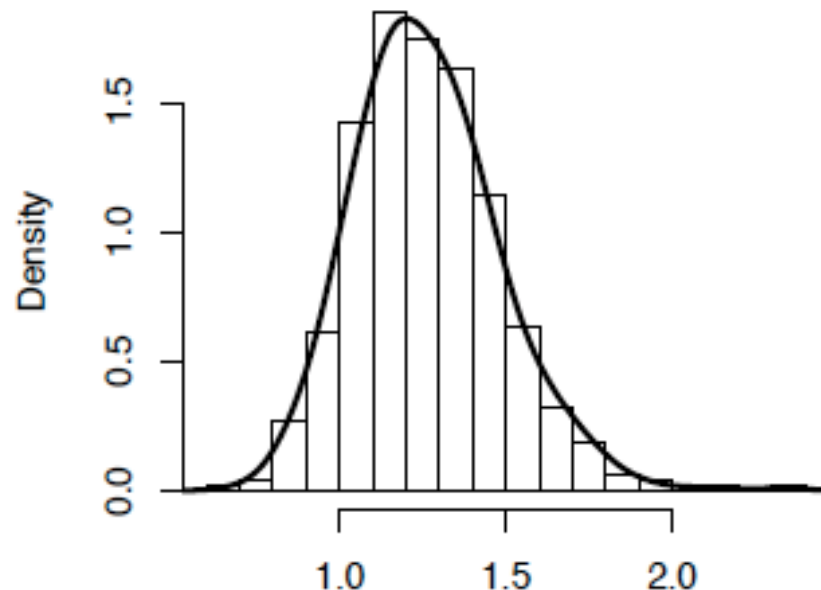
N. Thompson Hobbs





March 27, 2017



The first question we should ask after fitting a model: *Are the predictions of the model consistent with the data?*

- Is our deterministic model a reasonable representation of the process?
- Have we made the right choices of distributions to represent the uncertainties?



**A****B****C****D**

If we have a function  $f(A,B)$  specifying the joint probability of the continuous random variables  $A$  and  $B$ , then

$\int_B f(A,B)dB$  is the marginal probability of  $A$

and

$\int_A f(A,B)dA$  is the marginal probability of  $B$ .

This same idea applies to any number of jointly distributed random variables. We simply integrate over all but one.

# Posterior predictive checks

$$\left[ y^{new} \mid \mathbf{y} \right] = \underbrace{\int_{\theta_1} \dots \int_{\theta_n} [y^{new} \mid \theta_1 \dots \theta_n] [\theta_1 \dots \theta_n \mid \mathbf{y}] d\theta_1 \dots d\theta_n}_{\text{Posterior Predictive Distribution}}$$

It is called posterior because it is conditional on the observed  $\mathbf{y}$  and predictive because it is a prediction for an observable  $y^{new}$ . It gives the probability of a new prediction of  $y$  conditional on  $\boldsymbol{\theta}$ , which, in turn, is conditional on the data in hand,  $\mathbf{y}$ . Note that it is a marginal distribution because we are integrating over the  $\boldsymbol{\theta}$ .

$$\mu_i = g(\theta_1, \theta_2, \theta_3, x_i)$$

$$y_i \sim \text{normal}(\mu_i, \sigma^2)$$

Also see box 8.1 in  
Hobbs and Hooten

A new data set at each iteration

$k$	$\theta_1$	$\theta_1$	$\theta_3$	$i = 1$	$i = 2$	$i = 3$	$\dots$	$i = Y$
1	.42	3.3	20.3	$y_{1,1}^{new}$	$y_{1,2}^{new}$	$y_{1,3}^{new}$	$\dots$	$y_{1,Y}^{new}$
2	.41	2.3	18.5	$y_{2,1}^{new}$	$y_{1,2}^{new}$	$y_{1,3}^{new}$	$\dots$	$y_{1,Y}^{new}$
3	.46	3.1	16.6	$y_{3,1}^{new}$	$y_{1,2}^{new}$	$y_{1,3}^{new}$	$\dots$	$y_{1,Y}^{new}$
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$		$\vdots$
$K$	.39	3.4	22.1	$y_{n,1}^{new}$	$y_{n,2}^{new}$	$y_{n,3}^{new}$	$\dots$	$y_{1,Y}^{new}$

# This is easier done than said.

We have a model  $g(\theta, x)$  that predicts a response  $y$ .

We estimate the posterior distribution,  $[\theta | y]$ .

For any given value of  $x_i$ , we can simulate the posterior predictive distribution  $y_i^{new}$  by making a draw from  $[y_i^{new} | g(\theta, x_i), \sigma^2]$ . In MCMC, this simply means making draws from the data model at each iteration because each draw is conditional on the current values of the parameters. We simulate a new dataset by repeating these draws for all values of the  $x$ .

Accumulating many of these draws defines the posterior predictive distribution in exactly the same way that many draws allow us to define the posterior distribution of the parameters.



$$g(b_0, b_1, x_i) = b_0 + b_1 x_i$$

$$[b_0, b_1, \tau \mid \mathbf{y}] \propto \prod_{i=1}^n \text{normal}(y_i \mid g(b_0, b_1, x_i), \tau) \times$$

$$\text{normal}(b_0 \mid 0.0001) \text{normal}(b_1 \mid 0, .0001) \text{gamma}(\tau \mid .01, .01)$$

```

model{
  b0 ~ dnorm(0, .0001)
  b1 ~ dnorm(0, .0001)
  tau ~ dgamma(.01, .01)
  sigma <- 1/sqrt(tau)
  for(i in 1:length(y)){
    mu[i] <- b0 + b1*x[i]
    y[i] ~ dnorm(mu[i], tau)
    #posterior predictive distribution of y.new[i]
    y.new[i] ~ dnorm(mu[i], tau)
  }
}

```

# Posterior Predictive Checks

$T(\mathbf{y}, \theta)$  is a test statistic (e.g., mean, standard deviation, CV, quantile, or sums of squares discrepancy) calculated from the observed data.

$T(\mathbf{y}^{new}, \theta)$  is the corresponding statistic from the new "data" from the posterior predictive distribution.

We calculate:

$$P_B = \Pr\left(T(\mathbf{y}^{new}, \theta) \geq T(\mathbf{y}, \theta) \mid \mathbf{y}\right)$$

If  $P_B$  is very large or very small, then the difference between the observed data and the simulated data cannot be attributed to chance. This indicates lack of fit.

# Candidates for test statistics

- Mean
- variance
- Coefficient of variation
- quantiles
- maximum, minimum
- discrepancy: (observation - prediction)<sup>2</sup>
- chi-square:  $T(y, \theta) = \sum_i \frac{(y_i - E(y_i | \theta))^2}{\text{var}(y_i | \theta)}$
- deviance:  $T(y, \theta) = -2 \log[y | \theta]$

# R. A. Fischer's Ticks

A simple example: We want to know (for some reason) the average number of ticks on sheep. We round up 60 sheep and count ticks on each one. Does a Poisson distribution fit the distribution of the data?

$$[\lambda \mid \mathbf{y}] \propto \prod_{i=1}^{60} \text{Poisson}(y_i \mid \lambda) [\lambda]$$

For each value of  $\lambda$  in the MCMC chain, we generate a new data set,  $\mathbf{y}^{\text{new}}$ , by sampling from


$$y_i^{\text{new}} \sim \text{Poisson}(\lambda)$$

```
model{
lambda ~ dgamma(0.001,0.001)
for(i in 1:60){
  y[i] ~ dpois(lambda)
  y.new[i] ~ dpois(lambda) #simulate a new data set of 60 points
}
cv.y <- sd(y[ ])/mean(y[ ])
cv.y.new <- sd(y.new[ ])/mean(y.new[ ])
pvalue.cv <- step(cv.y.new-cv.y) # find Bayesian P value--the mean of
many 0's and 1's returned by the step function, one for each iteration in
the chain. The function step(z) returns a 1 if z > 0, returns 0
otherwise.
mean.y <-mean(y[ ])
mean.y.new <-mean(y.new[ ])
pvalue.mean <-step(mean.y.new - mean.y)
for(j in 1:60){
  sq[j] <- (y[j]-lambda)^2
  sq.new[j] <- (y.new[j]-lambda)^2
}
fit <- sum(sq[ ])
fit.new <- sum(sq.new[ ])
pvalue.fit <- step(fit.new-fit)
} #end of model
```

Key bit!




## Real Data




Simple model

## Simulated Data



# Posterior predictive check



```
}  
fit <- sum(sq[])  
fit.new <- sum(sq.new[])  
pvalue.fit <- step(fit.new-fit)  
} #end of model
```

## Simple model

P value for CV= .0013

P value for mean = .51

Remember, this is a two-tailed probability, so values close to 0 and 1 indicate lack of fit.

How could you modify this model to allow “extra” variance? Draw a Bayesian network and write out the posterior and joint distributions. You may not use the negative binomial. Hint-remember the owls fecundity model.



# Hierarchical model


```
model{
a~ dgamma(.001,.001)
b~ dgamma(.001,.001)
for(i in 1:60){
  lambda[i] ~ dgamma(a,b)
  y[i] ~ dpois(lambda[i])
  y.sim[i] ~ dpois(lambda[i])
}
cv.y <- sd(y[ ])/mean(y[ ])
cv.y.sim <- sd(y.sim[ ])/mean(y.sim[ ])
pvalue.cv <- step(cv.y.sim-cv.y) # find Bayesian P
value--the mean of many 0's and 1's returned by
the step function, one for each step in the chain
mean.y <-mean(y[])
mean.y.sim <-mean(y.sim[])
pvalue.mean <-step(mean.y.sim - mean.y)
for(j in 1:60){
  sq[j] <- (y[j]-lambda[j])^2
  sq.new[j] <- (y.sim[j]-lambda[j])^2
}
fit <- sum(sq[])
fit.new <- sum(sq.new[])
pvalue.fit <- step(fit.new-fit)
} #end of model
```

$$[a, b, \boldsymbol{\lambda} | \mathbf{y}] \propto \prod_{i=1}^{60} [y_i | \lambda_i] [\lambda_i | a, b] [a] [b]$$


Include **pvalue.fit** in variable names list for `coda.samples` or `jags.samples`. Report the mean of **pvalue.fit**



**Real Data**




**Simulated Data**



Hierarchical model

# Posterior predictive check



Hierarchical model

P value for CV = .45

P value for mean = .50